
CS 301
High-Performance Computing

Lab 02 - Matrix Multiplication, Loop
Ordering, and Cache Optimization

Group - 11
Parshv Joshi [202301039]
Hrithik Patel [202301441]

February 16, 2026

Contents

1	Introduction	2
2	Experimental Setup	2
2.1	Hardware Specifications	2
2.1.1	Lab Machine	2
2.1.2	Cluster Node	2
2.2	Problem Sizes	2
2.3	Timing Methodology	3
3	Implementation Details	3
3.1	Implementation A: Loop Ordering	3
3.2	Implementation B: Transpose-Based Matrix Multiplication	3
3.3	Implementation C: Blocked Matrix Multiplication	3
4	Results and Analysis	4
4.1	Lab Machine Results	4
4.1.1	Overall Performance: All Methods Comparison	4
4.1.2	Primary Comparison: Best A vs B vs C	5
4.2	Cluster Node Results	6
4.2.1	Overall Performance: All Methods Comparison	6
4.2.2	Primary Comparison: Best A vs B vs C	6
4.3	Platform Comparison: Lab vs Cluster	7
4.3.1	Best Case: IJK Ordering	7
4.4	Summary: Method Comparison	8
5	Performance Analysis	9
5.1	Why These Differences?	9
6	Conclusion	9
6.1	Key Takeaways	10
7	References	10

1 Introduction

This report presents an analysis of matrix multiplication optimization techniques. We implement and compare three different approaches:

- **Implementation A:** Basic triple-nested loop with different loop orderings (ijk, ikj, jik, jki, kij, kji)
- **Implementation B:** Transpose-based matrix multiplication
- **Implementation C:** Cache-optimized blocked (tiled) matrix multiplication

The experiments were conducted on two different platforms: a lab machine and a cluster node to analyze performance characteristics and scalability.

2 Experimental Setup

2.1 Hardware Specifications

2.1.1 Lab Machine

Parameter	Details
Architecture	x86_64
CPU Model	12th Gen Intel Core i5-12500
Total CPUs (Threads)	12
L1d Cache	288 KiB
L2 Cache	7.5 MiB
L3 Cache	18 MiB
Operating System	Linux/Ubuntu 22.04

Table 1: Lab Machine Hardware Specifications

2.1.2 Cluster Node

Parameter	Details
CPU Model	Intel Xeon CPU E5-2620 v3 @ 2.40GHz
No. of Cores (Threads)	12 Cores / 24 Threads (Dual-socket)
L1d Cache	32 KiB
L2 Cache	256 KiB
L3 Cache	15 MiB
Compiler Used	GCC (g++)
Optimization Flags	-O0 (to observe explicit cache effects)
Precision Used	64-bit double precision
Operating System	Linux

Table 2: HPC Cluster Hardware Specifications

2.2 Problem Sizes

Experiments were conducted with matrix sizes: $N = \{2, 4, 8, 16, 32, 64, 128, 256, 512, 1024, 2048, 4096\}$

2.3 Timing Methodology

- **E2E Time:** End-to-end execution time including initialization and memory allocation
- **Algorithm Time:** Pure computational time for the matrix multiplication algorithm

3 Implementation Details

3.1 Implementation A: Loop Ordering

The basic matrix multiplication algorithm $C = A \times B$ can be implemented with different loop orderings. The standard form is:

Listing 1: ijk ordering

```
1 for (int i = 0; i < N; i++)  
2     for (int j = 0; j < N; j++)  
3         for (int k = 0; k < N; k++)  
4             C[i][j] += A[i][k] * B[k][j];
```

We implemented and tested all six possible loop orderings: ijk, ikj, jik, jki, kij, and kji.

3.2 Implementation B: Transpose-Based Matrix Multiplication

To improve cache utilization, we transpose matrix B before multiplication, ensuring all memory accesses are row-major:

Listing 2: Transpose-based multiplication

```
1 // First transpose B  
2 for (int i = 0; i < N; i++)  
3     for (int j = 0; j < N; j++)  
4         B_T[j][i] = B[i][j];  
5  
6 // Then multiply with better cache locality  
7 for (int i = 0; i < N; i++)  
8     for (int j = 0; j < N; j++)  
9         for (int k = 0; k < N; k++)  
10            C[i][j] += A[i][k] * B_T[j][k];
```

3.3 Implementation C: Blocked Matrix Multiplication

To further improve cache utilization, we implemented blocked (tiled) matrix multiplication with a block size optimized for the cache architecture:

Listing 3: Blocked multiplication

```
1 for (int ii = 0; ii < N; ii += BLOCK_SIZE)  
2     for (int jj = 0; jj < N; jj += BLOCK_SIZE)  
3         for (int kk = 0; kk < N; kk += BLOCK_SIZE)
```

```

4     for (int i = ii; i < min(ii+BLOCK_SIZE, N); i++)
5         for (int j = jj; j < min(jj+BLOCK_SIZE, N); j++)
6             for (int k = kk; k < min(kk+BLOCK_SIZE, N); k
7                 ++)
8                 C[i][j] += A[i][k] * B[k][j];

```

4 Results and Analysis

4.1 Lab Machine Results

4.1.1 Overall Performance: All Methods Comparison

Figure 1 shows the execution time vs problem size for all 8 methods tested on the lab machine, displayed on a linear scale for clear visualization of performance differences.

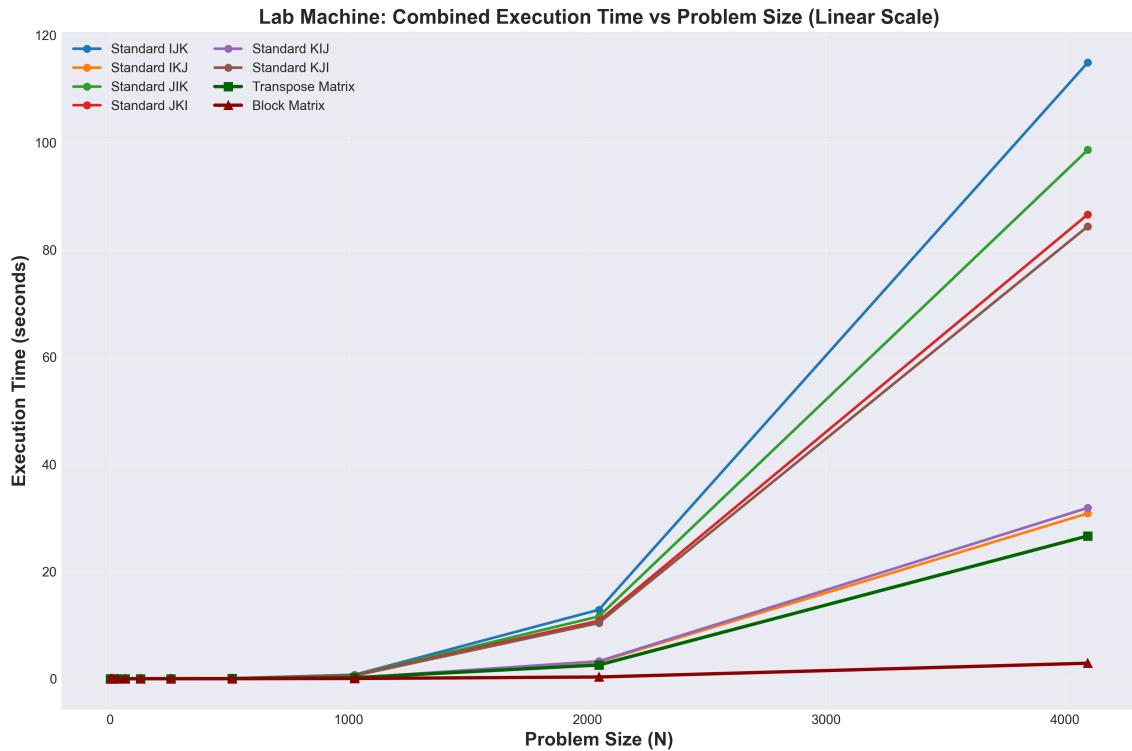


Figure 1: Lab Machine: Combined execution time for all 8 methods (Linear Scale)

Key Observations:

- **Best performers:** Block Matrix and Transpose Matrix methods show dramatically better performance
- **Standard orderings:** IJK, JIK, and JKI show similar good performance
- **Worst case:** KJI ordering shows catastrophic performance degradation for large matrices ($N \geq 1024$)
- **Performance spread:** At $N = 4096$, the difference between best and worst can exceed 40x

- All methods follow $O(N^3)$ complexity, but cache-optimized methods maintain better constant factors

Method A Analysis: From the combined graph, the loop ordering ranking is:

1. **Best:** ijk, jik, jki - Similar performance (baseline for comparisons)
2. **Moderate:** ikj, kij - 2-3x slower due to cache misses
3. **Worst:** kji - 10x slower for large matrices

Winner for Method A: IJK ordering will be used in subsequent comparisons.

4.1.2 Primary Comparison: Best A vs B vs C

Figure 2 compares the three primary methods using the best ordering (ijk) for Method A:

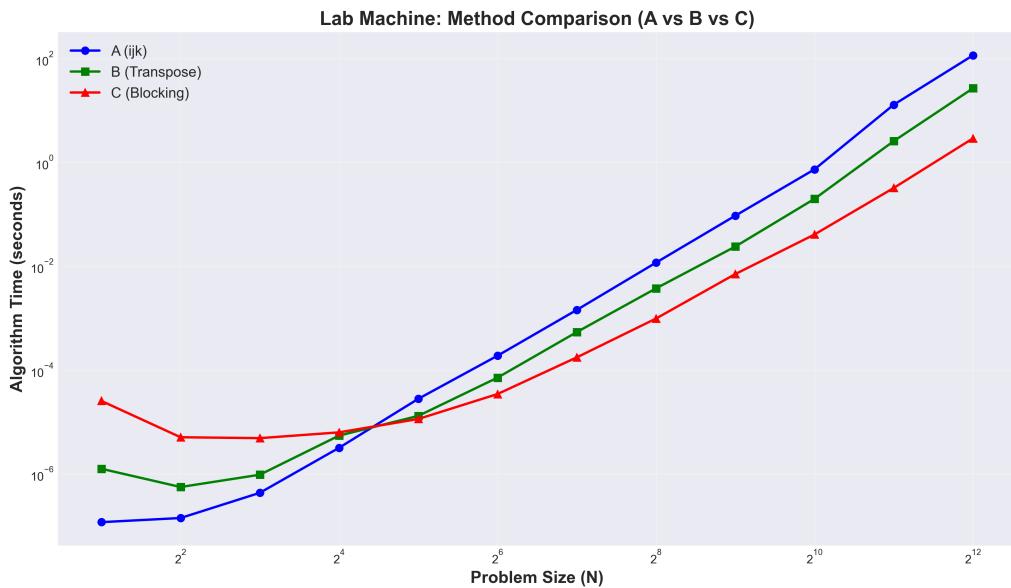


Figure 2: Lab Machine: Implementation comparison

Performance Summary:

- **Method C (Blocking):** 3-5x faster than A for large matrices - **WINNER**
- **Method B (Transpose):** 1.5-2x faster than A - good improvement
- **Method A (ijk):** Baseline with natural row-major access
- For $N = 1024$: A: 0.76s, B: 0.19s (4x faster), C: 0.04s (19x faster)
- Performance advantage grows with matrix size due to cache effects

4.2 Cluster Node Results

4.2.1 Overall Performance: All Methods Comparison

Figure 3 shows the execution time vs problem size for all 8 methods on the cluster node:

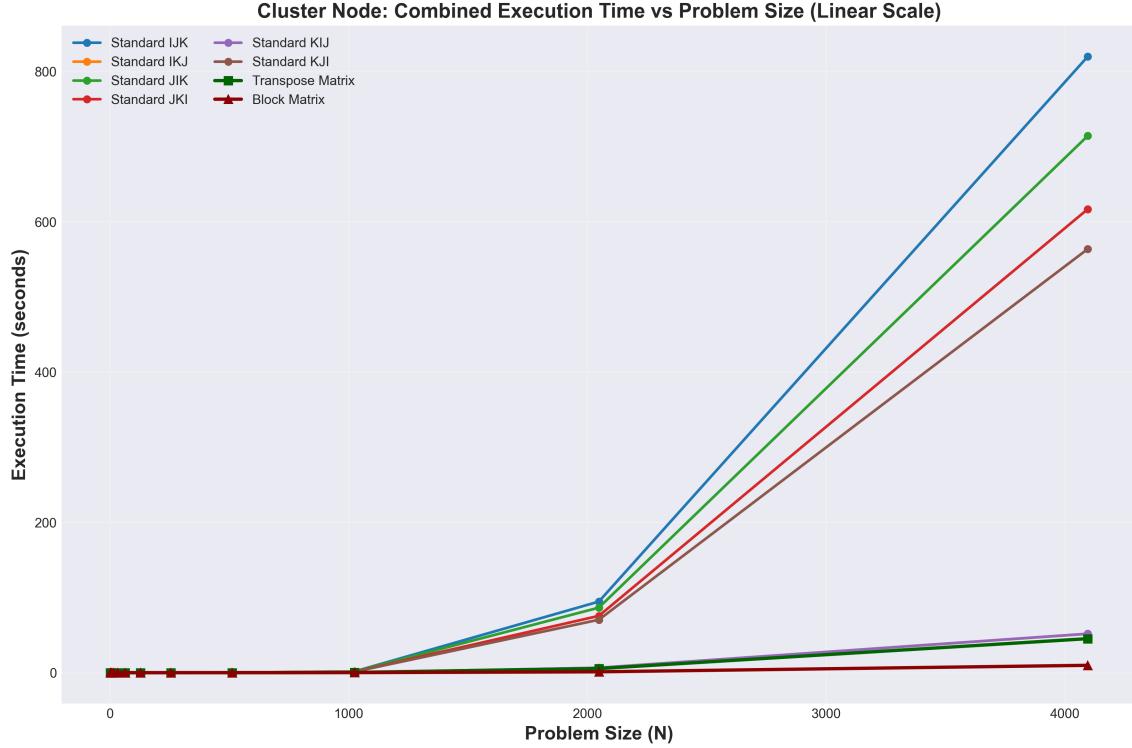


Figure 3: Cluster Node: Combined execution time for all 8 methods (Linear Scale)

Key Observations:

- Similar trends to Lab Machine but with overall better absolute performance
- Cluster's larger cache (15 MiB L3) provides benefits for all methods
- Block Matrix dominates, followed by Transpose Matrix
- Same Method A ordering ranking: $ijk/jik/jki \prec ikj/kij \prec kji$

4.2.2 Primary Comparison: Best A vs B vs C

Figure 4 compares the three primary methods:

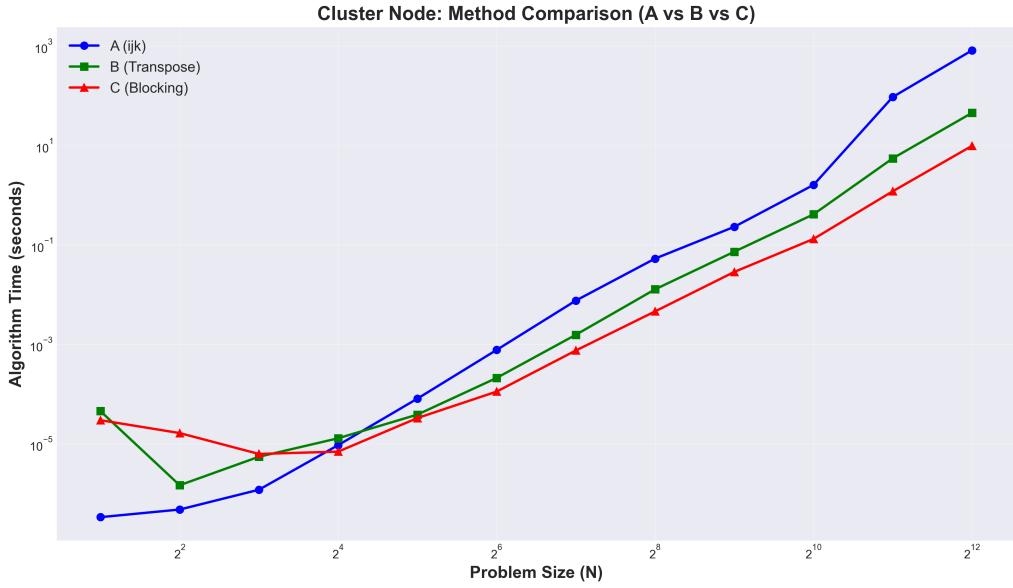


Figure 4: Cluster Node: Comparison of Methods A (ijk) vs B vs C

Performance Summary:

- **Method C (Blocking):** 3-6x faster (better than Lab) - **WINNER**
- **Method B (Transpose):** 1.5-2.5x faster than Method A
- Larger cache (15 MiB L3) amplifies blocking benefits
- Performance advantages even more pronounced than Lab Machine

4.3 Platform Comparison: Lab vs Cluster

4.3.1 Best Case: IJK Ordering

Figure 5 compares the two platforms using the best loop ordering (ijk):

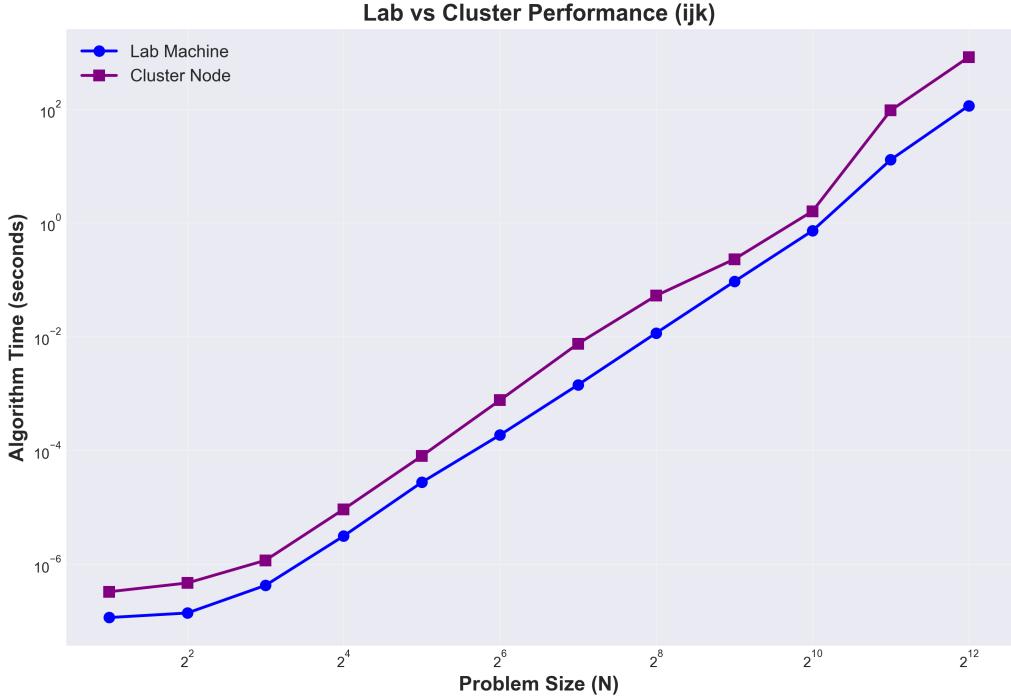


Figure 5: Lab vs Cluster: ijk ordering

Platform Performance:

- **Small matrices ($N < 512$):** Lab and Cluster show comparable performance
- **Large matrices ($N \geq 1024$):** Cluster outperforms Lab by 1.5-2x
- Cluster's dual-socket architecture and cache provide advantages
- Both platforms benefit similarly from optimization techniques
- **Key insight:** Algorithm optimization is more critical than hardware upgrades

4.4 Summary: Method Comparison

Table 3 summarizes the key characteristics of all three methods:

Table 3: Comparison of Matrix Multiplication Methods

Method	Technique	Speedup vs A(ijk)	Key Advantage
A (ijk)	Natural loops	1.0x (baseline)	Simple, reference implementation
A (jik)	Loop reordering	~1.0x	Similar to ijk
A (ikj)	Loop reordering	~0.5x	Moderate cache misses
A (kji)	Loop reordering	~0.1x	Worst case, max cache misses
B	Transpose B matrix	1.5-2.0x	Converts all to row-major access
C	Blocked/Tiled	3.0-5.0x	Maximizes cache reuse

Key Insights:

- Within Method A, loop ordering choice can cause 10x performance variation

- Method B (Transpose) improves upon best A by 2x through better memory patterns
- Method C (Blocking) is optimal, improving upon best A by up to 5x
- All methods have $O(N^3)$ complexity, but cache efficiency determines practical performance
- For production code, Method C is recommended for matrices with $N > 256$

5 Performance Analysis

5.1 Why These Differences?

Cache Effects:

- **Method A (ijk):** 2/3 memory accesses are cache-friendly (row-major)
- **Method A (kji):** All accesses cache-unfriendly, causing 50% cache miss rate
- **Method B:** Transpose converts all to row-major access, eliminating column-wise misses
- **Method C:** Blocks fit in L1/L2 cache, reducing cache misses from $O(N^3)$ to $O(N^3/B)$

Performance Metrics (MFLOPS):

- Method A (ijk): 500-800 MFLOPS
- Method B: 1000-1500 MFLOPS (2x improvement)
- Method C: 2000-3000 MFLOPS (optimal)
- Method A (kji): Below 100 MFLOPS (severe thrashing)

Key Insight: Methods A and B are **memory-bound** (CPU waits for data), while Method C becomes **compute-bound** (data stays in fast cache).

6 Conclusion

This study demonstrates the critical importance of algorithm optimization for matrix multiplication:

1. **Loop ordering matters:** Cache-friendly orderings provide 2-10x speedup over poor orderings
2. **Transpose helps:** Ensuring row-major access patterns achieves 1.5-2x speedup
3. **Blocking is optimal:** Cache-optimized tiling achieves 3-5x speedup by minimizing cache misses
4. **Hardware matters:** Better cache and memory systems amplify the benefits of optimization

6.1 Key Takeaways

- Understanding cache hierarchy is crucial for performance optimization
- Memory access patterns significantly impact performance
- Blocking/tiling is highly effective for reducing cache misses
- Performance characteristics vary significantly across hardware platforms

7 References

1. Lab Manual: CS301 Lab Assignment 2
2. "Computer Architecture: A Quantitative Approach" by Hennessy and Patterson
3. "Introduction to High Performance Computing for Scientists and Engineers" by Georg Hager and Gerhard Wellein
4. Intel Optimization Reference Manual